


Ant-v5 - penalised windows, all arms
penalty-block tail σ_{r+1}/σ_1



- baseline (stock SB3 TD3)
- FHR $\lambda=10, r=2$ (learned c)
- FHR $\lambda=10, r=8$ (learned c)
- FHR $\lambda=0.1, r=2$ (learned c)
- FHR $\lambda=0.1, r=8$ (learned c)
- FHR $\lambda=1, r=2$ (frozen c)
- FHR $\lambda=1, r=8$ (learned c)